



Module 7: Hands-On: Kubernetes Installation

Step 1: Launch 2 instances with the following configuration: **ubuntu 20.04 ami, t2.medium, sg: all traffic. ubuntu 20.04 ami, t2.micro , sg: all traffic**

Instance state = running X		Clear filters					
<input type="checkbox"/>	Name ▼	Instance ID	Instance state ▼	Instance type ▼	Status check	Alarm status	
<input type="checkbox"/>	K-worker	i-02c2fc5b4a8ce34c4	Running	t2.micro	2/2 checks passed	No alarms	+
<input type="checkbox"/>	K-master	i-0ecc3ba2e3d2ccb2f	Running	t2.medium	2/2 checks passed	No alarms	+

To Install Kubernetes use the following commands:

On Master and Worker node:

```
sudo su
```

```
apt-get update
```

```
apt-get install docker.io -y
```

```
service docker restart
```

```
curl -s https://packages.cloud.google.com/apt/doc/apt-key.gpg | apt-key add -
```

```
echo "deb http://apt.kubernetes.io/ kubernetes-xenial main" >/etc/apt/sources.list.d/kubernetes.list
```

```
apt-get update
```

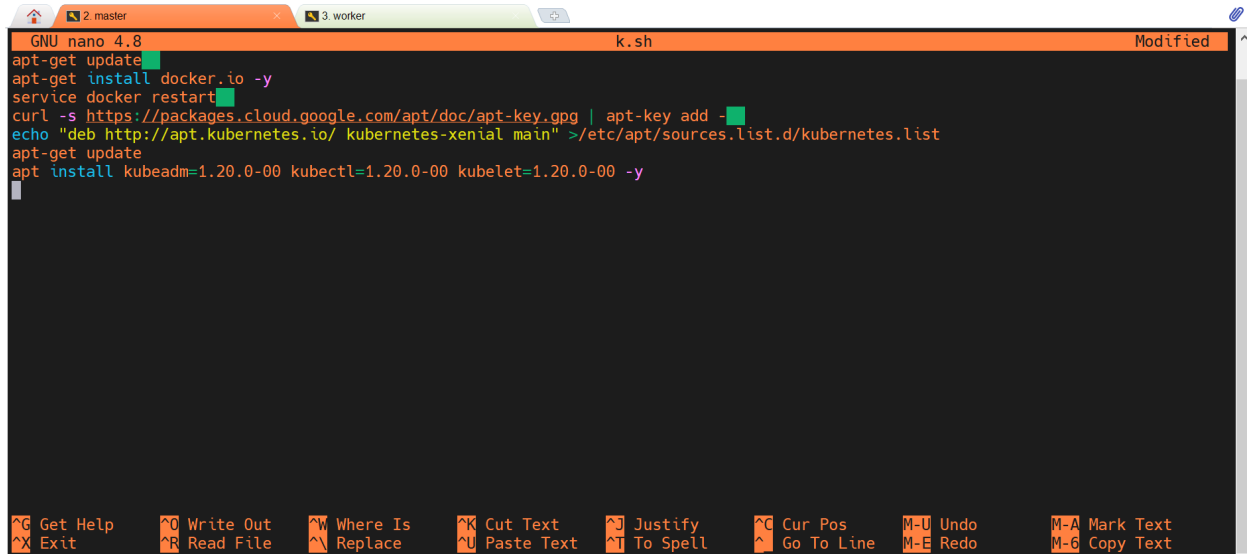
```
apt install kubeadm=1.20.0-00 kubectl=1.20.0-00 kubelet=1.20.0-00 -y
```

Step 2: On both master and worker nodes run the above command:

2.1. **sudo su**

2.2. create a script file **k.sh**

2.3. to execute the script file: **bash k.sh**

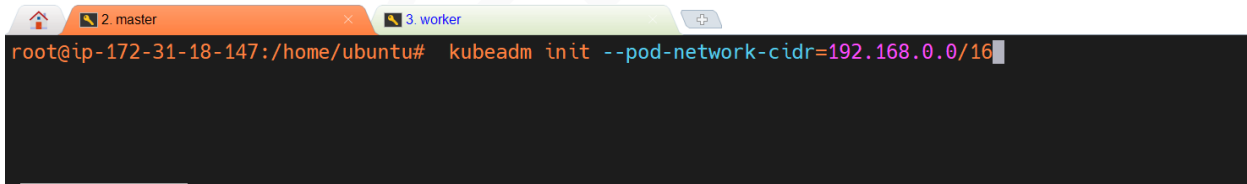


```
GNU nano 4.8 k.sh Modified
apt-get update
apt-get install docker.io -y
service docker restart
curl -s https://packages.cloud.google.com/apt/doc/apt-key.gpg | apt-key add -
echo "deb http://apt.kubernetes.io/ kubernetes-xenial main" >/etc/apt/sources.list.d/kubernetes.list
apt-get update
apt install kubeadm=1.20.0-00 kubectl=1.20.0-00 kubelet=1.20.0-00 -y
```

On Master:

Step 3:

kubeadm init --pod-network-cidr=192.168.0.0/16



```
root@ip-172-31-18-147:/home/ubuntu# kubeadm init --pod-network-cidr=192.168.0.0/16
```

```
[bootstrap-token] configured RBAC rules to allow the csrapprover controller automatically approve CSRs from a Node Bootstrap Token
[bootstrap-token] configured RBAC rules to allow certificate rotation for all node client certificates in the cluster
[bootstrap-token] Creating the "cluster-info" ConfigMap in the "kube-public" namespace
[kubelet-finalize] Updating "/etc/kubernetes/kubelet.conf" to point to a rotatable kubelet client certificate and key
[addons] Applied essential addon: CoreDNS
[addons] Applied essential addon: kube-proxy

Your Kubernetes control-plane has initialized successfully!

To start using your cluster, you need to run the following as a regular user:

mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config

Alternatively, if you are the root user, you can run:

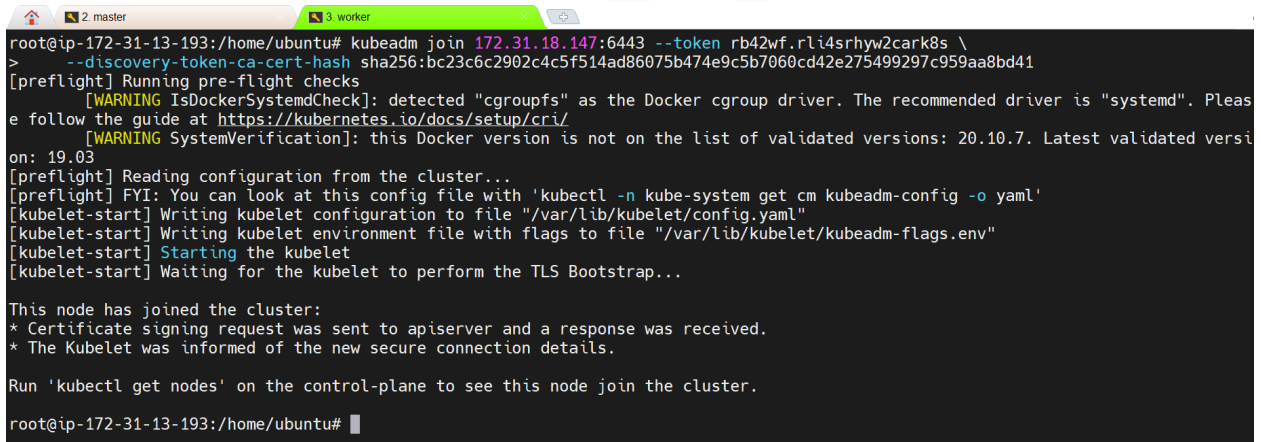
export KUBECONFIG=/etc/kubernetes/admin.conf

You should now deploy a pod network to the cluster.
Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:
https://kubernetes.io/docs/concepts/cluster-administration/addons/

Then you can join any number of worker nodes by running the following on each as root:

kubeadm join 172.31.18.147:6443 --token rb42wf.rli4srhyw2cark8s \
--discovery-token-ca-cert-hash sha256:bc23c6c2902c4c5f514ad86075b474e9c5b7060cd42e275499297c959aa8bd41
root@ip-172-31-18-147:/home/ubuntu#
```

Copy the token and paste it into the worker node



```
root@ip-172-31-13-193:/home/ubuntu# kubeadm join 172.31.18.147:6443 --token rb42wf.rli4srhyw2cark8s \
> --discovery-token-ca-cert-hash sha256:bc23c6c2902c4c5f514ad86075b474e9c5b7060cd42e275499297c959aa8bd41
[preflight] Running pre-flight checks
[WARNING IsDockerSystemdCheck]: detected "cgroups" as the Docker cgroup driver. The recommended driver is "systemd". Please
e follow the guide at https://kubernetes.io/docs/setup/cr
[WARNING SystemVerification]: this Docker version is not on the list of validated versions: 20.10.7. Latest validated versi
on: 19.03
[preflight] Reading configuration from the cluster...
[preflight] FYI: You can look at this config file with 'kubectl -n kube-system get cm kubeadm-config -o yaml'
[kubelet-start] Writing kubelet configuration to file "/var/lib/kubelet/config.yaml"
[kubelet-start] Writing kubelet environment file with flags to file "/var/lib/kubelet/kubeadm-flags.env"
[kubelet-start] Starting the kubelet
[kubelet-start] Waiting for the kubelet to perform the TLS Bootstrap...

This node has joined the cluster:
* Certificate signing request was sent to apiserer and a response was received.
* The Kubelet was informed of the new secure connection details.

Run 'kubectl get nodes' on the control-plane to see this node join the cluster.

root@ip-172-31-13-193:/home/ubuntu#
```

Step 4:

On Master:

```
exit
```

```
mkdir -p $HOME/.kube
```

```
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
```

```
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

Note: In case we want to retrieve the join token use the below-mentioned command.

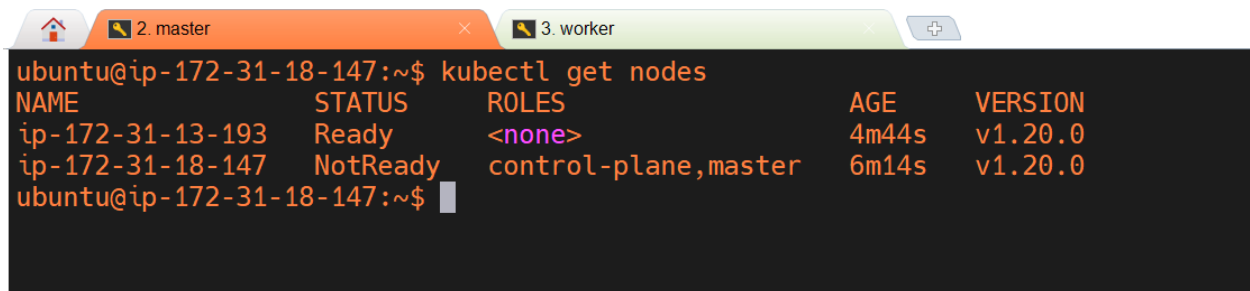
```
kubeadm token create --print-join-command
```

Step 5:

On Master:

```
kubectl apply -f https://docs.projectcalico.org/manifests/calico.yaml
```

```
kubectl apply -f  
https://raw.githubusercontent.com/kubernetes/ingress-nginx/controller-v0.49.0/deploy/static/provider/baremetal/deploy.yaml
```



```
ubuntu@ip-172-31-18-147:~$ kubectl get nodes
NAME                                STATUS    ROLES                  AGE      VERSION
ip-172-31-13-193                    Ready     <none>                 4m44s    v1.20.0
ip-172-31-18-147                    NotReady  control-plane,master   6m14s    v1.20.0
ubuntu@ip-172-31-18-147:~$
```

Our Kubernetes installation and configuration is complete.